

APIT: Concurrency

Simon Rogers

30/10/2014

Contents

Overview	1
What is concurrency?	1
Creating threads in Java	2
The options	2
Implementing the Runnable interface	2
Extending the Thread class	5
Thread names	6
Blocking methods	7
Interrupted Exception	7
Sleeping threads	7
Join	8
The benefits of parallel processing	8
Merge sort	8
Shared variables	8
Synchronized	12
Locks	12
Deadlocks	13
Conditions	13
Threads in Swing	14
The event dispatch thread	16
Longer jobs - the SwingWorker class	16
Other swing thread operations	17
Swing example - Game Of Life	17

Overview

- Concurrency
 - Threading
 - Solving threading problems
 - Threads in Swing
-

What is concurrency?

- Multiple parts of the program running simultaneously

- Why?
 - Make use of multiple processors
 - Efficient integration with slow devices (e.g. disks)
 - User-friendly-ness (responsive OS)
 - Useful (but oldish) paper
 - In Java, we can use **threads** to build concurrent programs
-

Creating threads in Java

The options

There are two ways of creating threads in Java:

- Implementing the **Runnable** interface
 - Extending the **Thread** class
-

Implementing the Runnable interface

The **Runnable** interface is very simple:

```
public interface Runnable {  
    public void run() {  
  
    }  
}
```

To use it:

- create a new class implementing this interface
 - create a **Thread** object
 - our new class *must* have a **run()** method.
-

For example, the following class implements **Runnable** and prints a particular **String** **n** times:

```
public class PointlessPrint implements Runnable {  
    private String message;  
    private int n;  
    public PointlessPrint(String message,int n) {  
        this.message = message;  
        this.n = n;  
    }  
    public void run() {  
        for(int i=0;i<n;i++) {  
            System.out.println(i + "/" +  
                n + " " + message);  
        }  
    }  
}
```

To use this class, we must create an instance and then place this instance within an instance of the `Thread` object:

```
public class RunnableTest {
    public static void main(String[] args) {
        PointlessPrint p = new PointlessPrint("Hello",100);
        Thread t = new Thread(p);
        t.start();
    }
}
```

`t.start()` starts the thread by invoking the `run()` method of `PointlessPrint`

From now on, we'll just use nested classes for things like this, to make the code a bit more concise:

```
public class RunnableTest {
    private static class PointlessPrint implements Runnable {
        private String message;
        private int n;
        public PointlessPrint(String message,int n) {
            this.message = message;
            this.n = n;
        }
        public void run() {
            for(int i=0;i<n;i++) {
                System.out.println(i + "/" + n + " " + message);
            }
        }
    }
    public static void main(String[] args) {
        PointlessPrint p = new PointlessPrint("Hello",100);
        Thread t = new Thread(p);
        t.start();
    }
}
```

Question: What would happen if `PointlessPrint` wasn't static?

```
public class RunnableTest2 {
    private class PointlessPrint implements Runnable {
        private String message;
        private int n;
        public PointlessPrint(String message,int n) {
            this.message = message;
            this.n = n;
        }
        public void run() {
            for(int i=0;i<n;i++) {
                System.out.println(i + "/" + n + " " + message);
            }
        }
    }
    public static void main(String[] args) {
        RunnableTest2 r = new RunnableTest2();
        PointlessPrint p = r.new
```

```

        PointlessPrint("Hello",100);
        Thread t = new Thread(p);
        t.start();
    }
}

```

The whole point of threads is that we can simultaneously create many of them. This is straightforward via an array of `Thread` objects:

```

public static void main(String[] args) {
    int nThreads = 2;
    Thread[] threads = new Thread[nThreads];
    for(int i=0;i<nThreads;i++)
    {
        PointlessPrint p = new PointlessPrint(
            "I am thread " + i,10);
        threads[i] = new Thread(p);
        threads[i].start();
    }
}

```

Producing the following output:

```

0/10 I am thread 0
1/10 I am thread 0
0/10 I am thread 1
1/10 I am thread 1
2/10 I am thread 0
2/10 I am thread 1
.
.
6/10 I am thread 1
7/10 I am thread 0
7/10 I am thread 1
8/10 I am thread 0
8/10 I am thread 1
9/10 I am thread 0
9/10 I am thread 1

```

We can see from the order of the `println` statements that both threads are running at the same time.

- The order will change every time we run it
 - The program stops once all threads are complete
 - It's impossible for us to know when Java switches from one thread to another
 - Note: They're not necessarily on different processors / cores, but might be
 - Note: Java always puts the main method in its own thread (you've been using threads all along!)
-

As an aside, here are two alternative `main` methods that would do the same thing. They should both make sense to you:

```

public static void main(String[] args) {
    int nThreads = 2;

```

```

Thread[] threads = new Thread[nThreads];
// Declare p outside the loop if we want future access to it
PointlessPrint[] p = new PointlessPrint[nThreads];
for(int i=0;i<nThreads;i++)
{
    p[i] = new PointlessPrint("I am thread " + i,10);
    threads[i] = new Thread(p[i]);
    threads[i].start();
}
}

public static void main(String[] args) {
    int nThreads = 2;
    Thread[] threads = new Thread[nThreads];
    for(int i=0;i<nThreads;i++)
    {
        // Anonymous PointlessPrint object
        threads[i] = new Thread(new PointlessPrint("I am thread " + i,10));
        threads[i].start();
    }
}

```

Extending the Thread class

- The alternative to implementing the `Runnable` interface
 - Create a new class that extends `Thread`
 - The new class has to have a method that overrides `run()`
 - The equivalent to our previous example can be found in `SimpleThreadTest`
-

```

public class SimpleThreadTest {
    private static class PointlessPrint extends Thread {
        private String message;
        private int n;
        public PointlessPrint(String message, int n) {
            this.message = message;
            this.n = n;
        }
        public void run() {
            for(int i=0;i<n;i++) {
                System.out.println(i + "/" + n + " " + message);
            }
        }
    }
}

public static void main(String[] args) {
    PointlessPrint[] threads = new PointlessPrint[2];
    for(int i=0;i<2;i++) {
        threads[i] = new PointlessPrint("I am thread " + i,10);
        threads[i].start();
    }
}
}

```

Thread names

- In our examples, we passed a message to a thread to help identify it
- Threads can also be given names through their constructor:

```
Thread t = new Thread(aRunnableThing, "my name");
Thread t = new Thread("my name");
```

- which can be accessed via:

```
thread.getName()
```

- See notes for examples

Example when implementing Runnable:

```
public class ThreadNameTest {
    private static class PointlessThread implements Runnable {
        private int n;
        public PointlessThread(int n) {
            this.n = n;
        }
        public void run() {
            for(int i=0; i<n; i++) {
                System.out.println(Thread.currentThread().getName() + " " + i);
            }
        }
    }
    public static void main(String[] args) {
        Thread[] threads = new Thread[2];
        for(int i=0; i<2; i++) {
            threads[i] = new Thread(new PointlessThread(10), "I am " + i);
            threads[i].start();
        }
    }
}
```

Example when extending Thread:

```
public class ThreadNameTest2 {
    private static class PointlessThread extends Thread {
        private int n;
        public PointlessThread(int n, String name) {
            super(name); // Thread constructor
            this.n = n;
        }
        public void run() {
            for(int i=0; i<n; i++) {
                System.out.println(this.getName() + " " + i);
            }
        }
    }
    public static void main(String[] args) {
        PointlessThread[] threads = new PointlessThread[2];
        for(int i=0; i<2; i++) {
            threads[i] = new PointlessThread(10, "Thread " + i);
            threads[i].start();
        }
    }
}
```

```
}  
}
```

Blocking methods

- Blocking methods are methods that rely on something else within the system for termination
 - Waiting for a timer to elapse
 - Waiting for another thread to end
 - Because these methods rely on something external, they might be waiting forever.
 - To ensure smooth running, they should be *cancelable*
-

Interrupted Exception

- Threads can be interrupted by other threads
 - When a thread is interrupted, one of two things happen:
 - If it is running an interruptable method (e.g. `Thread.sleep()`), the method unblocks and throws the `InterruptedException`
 - Otherwise, its (boolean) interrupted status is set
 - Interrupted status can be read with `Thread.isInterrupted()`
 - Interrupted status can be read and reset (0) with `Thread.interrupted()`
-

The following code will start a thread, wait a random amount of time and then interrupt it. When it is interrupted, it just finishes.

```
import java.util.Random;  
public class InterruptTest {  
    public static class InterruptableThread implements Runnable {  
        public void run() {  
            int i=0;  
            while(Thread.currentThread().isInterrupted()==false) {  
                System.out.println(i++);  
            }  
        }  
    }  
    public static void main(String[] args) {  
        Thread t = new Thread(new InterruptableThread());  
        t.start();  
        int r = new Random().nextInt();  
        for(int i=0;i<r;i++) {  
            // Nothing, just eat up some time  
        }  
        t.interrupt();  
    }  
}
```

Sleeping threads

- `Thread.sleep(long time)` is a method that throws `InterruptedException` so it has to be caught:

```

public void run() {
    try {
        Thread.sleep(1000);
    } catch (InterruptedException e) {
        System.out.println("You woke me up");
    }
}

```

Join

- In many applications it will be useful to know if a Thread has finished.
 - `someThread.join()` pauses the current thread until `someThread` has finished.
 - `join()` throws an `InterruptedException`
-

The benefits of parallel processing

- Many machines have multiple cores / processors
 - Threads can be placed on different cores / processors
 - Java does this for us - we have no control
 - Running things in parallel should give us speed improvements
 - Although it increases system book-keeping
 - Class example: merge sort
-

Merge sort

- We'd like to sort the values in a large array.
 - Can be made parallel:
 - Split the array into N smaller arrays
 - Sort the smaller arrays
 - Merge the results together
 - How much speed-up will this give?
-

Shared variables

- A benefit of threads is the shared address space
 - Multiple threads can access the same shared resources
- For example, suppose I would like to make a system where several threads can all increment the same counter
- See `CounterExample`
- Why can't we just pass an `Integer` around instead of a `MyCounter` object? (see Immutable objects and Immutable objects 2)

```

public class CounterExample {
    public static class MyCounter {
        // We need this method because ints are immutable

```



```

    private int count = 0;
    public int getCount() {
        return count;
    }
    public void setCount(int count) {
        this.count = count;
    }
}
public static class Counter extends Thread {
    private MyCounter count;
    private int n;
    public Counter(MyCounter count, int n) {
        this.count = count;
        this.n = n;
    }
    public void run() {
        for(int i=0; i<n; i++) {
            int temp = count.getCount();
            temp++;
            count.setCount(temp);
        }
    }
}
public static void main(String[] args) {
    MyCounter count = new MyCounter();
    Counter c = new Counter(count, 100);
    c.start();
    try {
        c.join();
    } catch (InterruptedException e) {
        //Do nothing
    }
    System.out.println(count.getCount());
}
}

```

This code outputs 100 at the end, as you might expect.

Question: try removing the join code (i.e. everything inside the try catch block) - what happens?

Let's now update main so that it creates several threads, all accessing the same MyCounter object (CounterExample2):

```

public static void main(String[] args) {
    MyCounter count = new MyCounter();
    int nCounters = 100;
    Counter[] c = new Counter[nCounters];
    for(int i=0; i<nCounters; i++) {
        c[i] = new Counter(count, 1000);
        c[i].start();
    }
    try {
        for(int i=0; i<nCounters; i++) {
            c[i].join();
        }
    } catch (InterruptedException e) {
    }
}

```

```

        //Do nothing
    }
    System.out.println(count.getCount());
}

```

If all works correctly, we should see 100 times 1000 (=100000). But we don't. Each time we run, we see something different (try it), always below 100000. What's happening?

```

public void run() {
    for(int i=0;i<n;i++) {
        int temp = count.getCount();
        temp++;
        count.setCount(temp);
    }
}

```

The problem is found in the `run()` method:

- Remember that we have no idea when Java will move from one Thread to another.
- If it moves between the `getCount()` and `setCount()` methods...

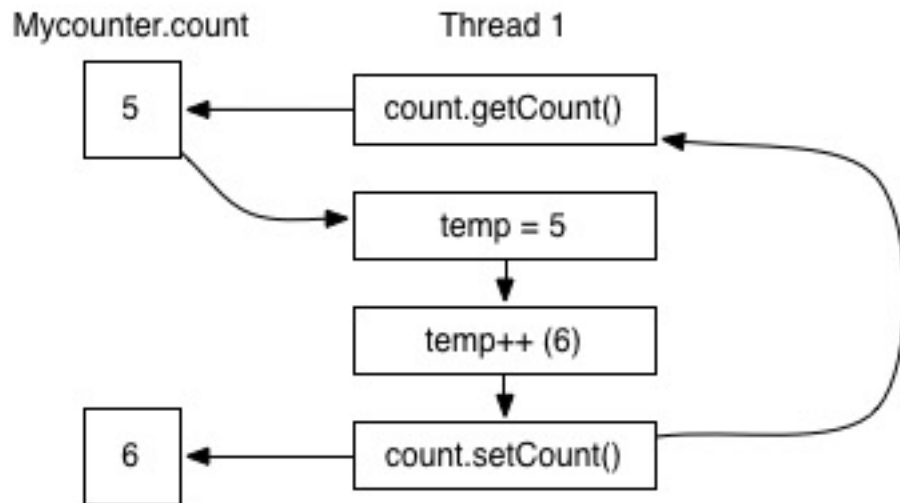


Figure 1: Single thread operation

-
- If control is passed between `get` and `set` the new thread sees an out of date value.
 - Remember: `temp` is local to each thread.
 - In reality, thread 1 might be sitting dormant with `temp=5` for a long time
 - When it finally updates it, it will effectively delete lots of updates performed by other threads
 - It is known as a *race condition*
 - This is a *big* problem in multi-threaded programs
 - We'll now look at ways of overcoming it
-

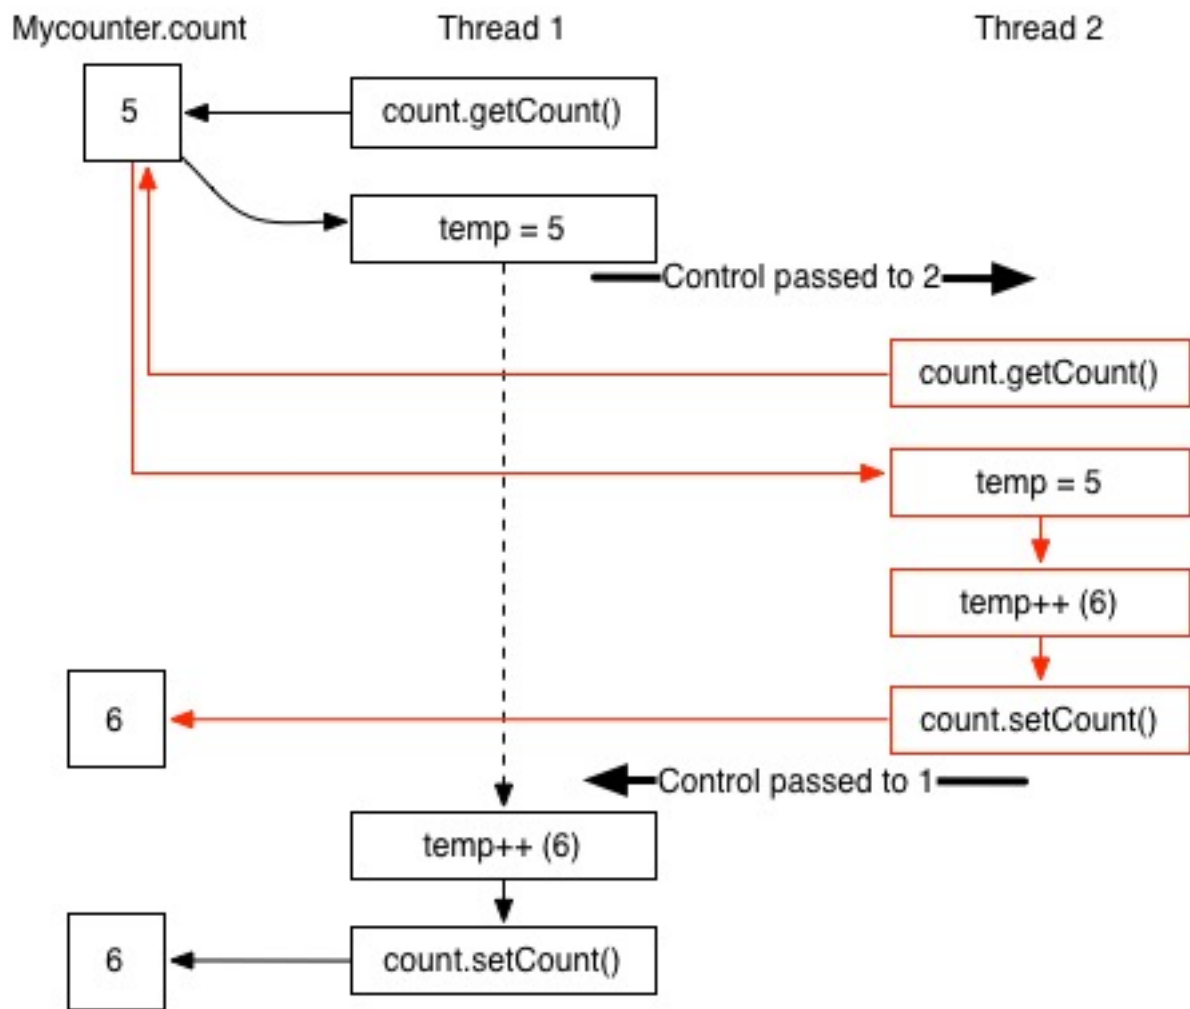


Figure 2: Multi-thread operation

Synchronized

- To overcome race conditions, we must *lock* objects.
 - All objects have an associated monitor that can be locked or unlocked (we don't see this)
 - A thread can lock a monitor, ensuring no other threads can modify it
 - Other threads trying are blocked until the lock is released
 - The easiest way is with `synchronized` blocks and methods.
-

- `CounterExample3` gives a new version of our Counter program
 - Note that the incrementation is now done inside the class
 - It still has the same problem, although perhaps not as extreme (try it and see. Why? Think about how `count++` is done and the chances of being interrupted at a bad point)
 - We can solve our race condition by making the `increment()` a synchronised method
 - when any thread is invoking a synchronized method, all other threads trying to are paused until it has finished
 - See `CounterExample4` ... problem solved
-

- Alternatively, we can just synchronize a block of code.
- E.g. instead of declaring `increment()` as `synchronized` we can:

```
public void run() {
    for(int i=0;i<n;i++) {
        synchronized(count) {
            count.increment()
        }
    }
}
```

- This causes the thread to lock the `count` object
 - No other threads can modify `count` when the thread is in this block.
 - This will also fix the problem - try it
 - There are other blocks that could be synchronized - try some
-

Locks

- An alternative approach involves creating `Lock` objects
- For example, `ReentrantLock()` (`CounterExample5`):

```
public static class MyCounter {
    private int count = 0;
    private ReentrantLock counterLock =
        new ReentrantLock();
    public void increment() {
        counterLock.lock();
        count ++;
        counterLock.unlock();
    }
    ...
}
```

- When `counterLock` is locked, no other thread can lock it until it has been unlocked (Phonebooth analogy in `BigJava`)

-
- There's a problem: if the code between `lock` and `unlock` throws an exception the `unlock` never happens
 - Phonebooth user collapsing??
 - Always do the following to ensure the lock is released:

```
someLock.lock();
try {
    // Some code
}
finally {
    someLock.unlock();
}
```

Deadlocks

- What if two threads are both waiting for one another to release a lock?
- The program will hang indefinitely
- This is a *deadlock*
- For example, suppose adding another object to our `CounterExample` that decrements `MyCounter`
 - If we set the system up so that in total the same amount is incremented and decremented then `count` will sometimes become negative (depending on ordering of events)
 - See `CounterDecounter`

-
- We would like to ensure this number never goes negative
 - One way of doing this would be to put some kind of wait condition in the `decrement` method (`CounterDecounter2`):

```
counterLock.lock();
try {
    while(count < amount) {
        Thread.sleep(1);
    } catch (InterruptedException e) {
        // fall through
    } finally {
        counterLock.unlock();
    }
}
```

- This causes the program to hang whenever it tries to decrement by an `amount` that is greater than `count`
 - Because the thread has locked `counterLock` no other thread can increase `amount`
 - This is a *deadlock*
-

Conditions

- Conditions allow threads to temporarily unlock locks whilst they await some condition to be fulfilled
- In this case, we'd like to temporarily unlock within a thread that is waiting to `decrement`
- Conditions are created from locks
- We can add a condition to `MyCounter` as follows:

```
private ReentrantLock counterLock = new ReentrantLock();
private Condition bigEnough = counterLock.newCondition();
```

- Threads can await the condition through the `Condition.await()` method
- We add this to our `decrement` method:

```
public void decrement(int amount) {
    counterLock.lock();
    try {
        while(count < amount) {
            bigEnough.await();
        }
        count -= amount;
        System.out.println("Subtracting " + amount + ", result " + count);
    } catch (InterruptedException e) {
        // Fall through
    } finally {
        counterLock.unlock();
    }
}
```

- A thread calling `decrement` when `count < amount` will wait until another thread invokes the `Condition.signalAll()` method.
- We put this method into the `increment` method:

```
public void increment(int amount) {
    counterLock.lock();
    try {
        count += amount;
        System.out.println("Adding " + amount + ", result " + count);
        bigEnough.signalAll();
    } finally {
        counterLock.unlock();
    }
}
```

- Whenever an increment is made, all threads waiting on this condition are restarted.
 - Note that the `signalAll()` method doesn't mean that `amount` is big enough
 - The syntax in `decrement()` means that `signalAll()` will cause the thread to check again.
 - It might just end up invoking `await()` again.
 - Run `CounterDecounter3` and verify that `count` never becomes negative
-

Threads in Swing

- In general Swing is not thread safe
 - You can't use normal threads
 - Ignore everything up until now!
- But, Swing does give you threading capabilities
- First, why do we need threads in Swing?

```

    - SwingThread

import javax.swing.JFrame;
import javax.swing.JButton;
import java.awt.event.ActionListener;
import java.awt.event.ActionEvent;
import javax.swing.JTextField;
import java.awt.GridBagLayout;

public class SwingThread extends JFrame implements ActionListener {
    private final JButton startButton, stopButton;
    private final JTextField countField, outField;
    public SwingThread() {
        super("Swing Thread");
        setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
        getContentPane().setLayout(new GridBagLayout());

        startButton = makeButton("Start");
        stopButton = makeButton("Stop");

        countField = makeText();
        outField = makeText();
        pack();
        setVisible(true);
    }
    private JTextField makeText() {
        JTextField b = new JTextField(30);
        getContentPane().add(b);
        return b;
    }
    private JButton makeButton(String caption) {
        JButton b = new JButton(caption);
        b.setActionCommand(caption);
        b.addActionListener(this);
        getContentPane().add(b);
        return b;
    }

    public void actionPerformed(ActionEvent e) {
        if(e.getActionCommand() == "Start") {
            outField.setText("You pressed start");
            startCounting();
        } else if(e.getActionCommand() == "Stop") {
            outField.setText("You pressed stop");
        }
    }
    private void startCounting() {
        try {
            for(int i=0; i<100; i++) {
                countField.setText(String.format("%d", i));
                Thread.sleep(100);
            }
        } catch (InterruptedException e) {
    
```

```

        }
    }
    public static void main(String[] args) {
        new SwingThread();
    }
}

```

When this code is run, the system waits patiently until `start` is pressed. When this happens, control is passed to the `actionPerformed()` method. This method sets the text in `outField` and then starts the counter. During counting, the code updates `countField` and sleeps a lot.

You might guess that on pressing `start` you would see the text in `outField` immediately and then see the text in `countField` gradually increasing. Unfortunately, this is not what happens. When we call `JTextField.setText()` we are changing the text that is stored. For the screen to update, Java must do various other processes (hidden from us). These can't be done if the system is busy. In this case, the system is busy counting and so nothing changes on screen until the counting has finished, when we see the final value and 'you pressed start'.

To overcome this, we need to use the thread objects that swing provides.

-
- System becomes unresponsive whilst counting
 - Nothing updates until counting has finished
 - until we exit the `actionPerformed` method
 - We need threads
-

The event dispatch thread

- Event handling code in Swing runs on the event dispatch thread
 - e.g. `actionListeners`
 - Things on this thread should be *short tasks*
 - otherwise system becomes unresponsive
 - Note: some swing component methods can be invoked from any thread (marked as *thread safe* in API)
 - Why isn't all of swing thread safe? read this
-

In our previous code, we unwittingly put the counter on the event dispatch thread and the system became unresponsive. Only very fast things should go on this thread (e.g. starting other threads that can run in the background).

Longer jobs - the `SwingWorker` class

- Long tasks should not be run on the event dispatch thread
 - Instead we use worker threads
 - Created by extending `SwingWorker`
 - The new class must extend:
 - `doInBackground()`
 - And can also use:
 - `publish()` and `process()` to display interim results
 - `done()` to invoke a method on the event dispatch thread when the task is complete.
 - `Counter.java`
-

- Note that `SwingWorker` takes two types:
 - `SwingWorker<A,B>`
 - A: the return value
 - B: the object passed by `publish`
- Note also that `process(List b)` takes a list
 - There may be many calls to `publish` before `process` is called
- The various Swing layout things should be things you've seen before?

`SwingWorker` objects allow us to run time consuming processes in the background without the system becoming unresponsive. When you extend the `SwingWorker` class, you *have* to implement `doInBackground()` which is the method called when we invoke the `execute()` method of the object. `doInBackground()` requires two object types to be specified: the first is its return value, the second is the type of object used by `publish()`. Other methods can also be implemented:

- `publish()`: this allows us to display preliminary results as our background task is running. Objects passed to `publish` can be retrieved when the system invokes `process()` ...
- `process()`: If we use `publish()` in `doInBackground()`, the system will periodically invoke `process()`. `Process` is passed a `List` of the objects passed to `publish()`. Often we'll only be interested in the last one. If you wanted to say update a `JTextArea` with current status, you should do so in `process()`.
- `done()`: This method is invoked on the event dispatch thread when `doInBackground()` has finished – useful for changing the enabled status of buttons etc, or displaying a message. Remember though, it's on the event dispatch thread so shouldn't be used to long tasks.

Other swing thread operations

- Initial threads (in `SwingUtilities`):
 - `invokeLater(Runnable go)` runs `go` on the event dispatch thread.
 - `invokeAndWait(Runnable go)` runs `go` on the event dispatch thread and then waits for it to finish.
 - Typically, these are used for starting the GUI (i.e. creating a `JFrame` object):

```
public static void main(String[] args) {
    SwingUtilities.invokeLater(new Runnable() {
        new SwingThread();
    });
}
```

Initial threads are typically just used for initially creating GUI objects. They put runnable objects onto the event dispatch thread. It is good practise to always create your initial objects in this manner.

Swing example - Game Of Life

- Class exercise: building a Game of Life simulator
- Details: Conway's Game of life
- We need a responsive application that animates a 'world' and allows users to start, stop, toggle cells, change speed, clear the world and randomise the world